

An Executable Safety Ladder Toward Unbounded Viewstamped Replication

Simon Massey

Abstract—Unbounded Viewstamped Replication Revisited (uVRR) seeks to combine replicated volatile state, low-delay primary replacement, and reconfiguration without imposing a fixed pipeline bound. These requirements cross several distinct proof boundaries. We audit an inherited Lean development of Turner’s era-indexed Paxos agreement argument and extend it with a finite-support weighted-majority theorem, a derivation of the proposal-era bound from suffix promises, operational acceptor and normal-log invariants, a whole-log selection induction step, and fault-injection controls. Every positive result is checked by Lean 4.33.1 using its core libraries. Executable Showboat documents preserve source, build commands, and theorem axiom reports. All eight inherited transcripts reproduce, although the eighth records an unsuccessful proof attempt. We distinguish this reproducibility result from mathematical correctness and from protocol refinement. We also separate counterexamples establishing independence of assumptions from claims of universal necessity. A section-by-section examination of Viewstamped Replication Revisited identifies the outstanding integration obligations: concrete protocol refinement, recovery fencing, repeated reconfiguration, client semantics, and progress. The classic crash-recover (CR) diskless-recovery amnesia class of Michael et al. (DISC’17) is eliminated by construction: uVRR evicts the crashed identity via Crash-Stop-Self-Evict reincarnation (a dirty node reopens under a new incarnation and rejoins as a weight-0 learner). We check the crash-vector collection primitive that literature proposes, without claiming a repair. Bounded Leanstrat experiments are reported as part of the methodology, including unsuccessful runs and unmeasured costs. This checkpoint establishes reusable verified components; an end-to-end uVRR correctness theorem and a millisecond failover measurement remain open.

Index Terms—Distributed algorithms, replication, reconfiguration, crash recovery, formal verification.

I. PROBLEM AND CONTRIBUTION

THE motivating uVRR proposal asks whether a strongly consistent replicated service can replace failed members while continuing to process requests, without forcing its log to disk on the normal agreement path [5]. A precise answer requires more than an intersection calculation. The membership used by a vote must be agreed; a replacement primary must preserve committed operations; a crashed replica must not cast votes using forgotten state; and a reply must correspond to an operation with the promised client semantics. Progress and elapsed time impose additional assumptions.

Turner’s UPaxos construction permits a sequence of consensus instances to remain unbounded during suitable reconfigurations [3]. Its safety argument provides a natural starting point. The inherited development encoded the Synod history

invariants, reduced an era-indexed history to them, checked two finite reconfiguration schedules, and isolated a leader’s casting-vote schedule. It explicitly excluded a proof of VRR itself. Our audit confirms the compiled mathematical results while narrowing several statements in the associated report.

The checked extensions begin with five positive results. First, an arbitrary finite-support version of UPaxos’s weighted intersection lemma replaces dependence on a four-node example. Second, the omitted suffix-promise and known-instance frontier conditions now derive the proposal-era inequality used in the agreement reduction. Third, induction over an acceptor transition system establishes the historical promise conditions for that component. Fourth, an executable whole-log selector is proved maximal, with a quorum-based prefix preservation induction step and an exact bounded-rank correspondence to the TLA model. Fifth, a fixed-view message induction derives same-view report comparability and replica prefix retention. Constructive negative controls test the role of selected assumptions. The laboratory book records intermediate failures, tool outcomes, and restart instructions.

The shared crash-free model additionally proves compatibility of committed logs across unbounded views. Recovery investigation produces a conditional fence-retention theorem; uVRR then removes the classic crash-recover class by construction, via Crash-Stop-Self-Evict reincarnation: a node whose durable superblocks record an unflushed session reopens under a new incarnation, and the leader evicts the old identity through the forced weight sequence (exiting 1 to 0, joining at 0, then 0 to 1). Classic VRR-2012 diskless quorum recovery (its section 4.3) is not performed, so the DISC’17 Appendix B.1 amnesia class is avoided rather than repaired.

The central unresolved claim is stated explicitly: *every execution of the intended uVRR message protocol, under a specified crash-recovery model, implements a linearizable replicated state machine across arbitrarily many reconfigurations*. The results below discharge parts of this claim. They do not discharge it by renaming a Paxos ballot a VRR view.

II. SETTING AND TRUST BOUNDARY

Let A , B , and V be node, ballot, and value types. A node set is a predicate $A \rightarrow \text{Prop}$; a quorum family is a predicate on such sets. Define

$$Q \curvearrowright R \iff \forall q \in Q, r \in R, \quad q \cap r \neq \emptyset.$$

This relation is symmetric. It is a property of two families, not a requirement that every pair of members within either family

intersect. Empty families make some such statements vacuous. Consequently nonemptiness conditions in the agreement theorem must remain visible.

The history model records messages that have ever been emitted. Its predicates include free promises $p(a, b)$, promises carrying a last vote $p(a, b; c)$, proposals, acceptances, and choices. A function $v(b)$ gives the unique value associated with each ballot. The representation presupposes ballot ownership or an equivalent mechanism preventing two values at one ballot; it is not a proof that a wire protocol enforces that mechanism.

The ballot order is well founded and strictly total. The era development uses $B = \mathbb{N} \times \mathbb{N}$ ordered lexicographically, so that unboundedly many rounds are available in each era. This is a mathematical domain; a fixed-width implementation needs a separate overflow argument.

The executable library imports Lean core and other modules in the same development. It has no Mathlib dependency. A proof may use standard Lean axioms `propext`, `Quot.sound`, and `Classical.choice`; these are distinguished from additional protocol hypotheses appearing in theorem types. Printing axioms alone does not expose every protocol assumption: a conditional theorem can have an empty axiom footprint and strong hypotheses. The central generic Synod theorem has no axiom dependencies. Other results use subsets of the three standard axioms.

III. CONDITIONAL AGREEMENT AND RECONFIGURATION

A. The Synod contract

The inherited structure `Synod.Inv` expresses six conditions:

- 1) If $b_2 < b_1$, b_1 is proposed, and b_2 is chosen, then $\mathcal{Q}^I(b_1) \cap \mathcal{Q}^II(b_2)$.
- 2) A free promise at b excludes every acceptance below b .
- 3) A report $p(a, b; c)$ names an accepted ballot $c < b$ and excludes acceptances strictly between c and b .
- 4) A proposal has a phase-I quorum of replies. If any report carries a vote, the proposed value equals that of a maximal reported ballot.
- 5) Every acceptance has a proposal.
- 6) Every choice has a phase-II quorum of acceptances.

The maximal reported ballot in condition 4 is supplied existentially. An implementation must obtain it from a finite collection and enforce the value rule. The history theorem does not construct a leader's message handler.

Theorem 1 (Synod agreement). *Assume the six conditions, a well-founded strict total ballot order, and nonempty phase-II quorums. If b and c are chosen, then $v(b) = v(c)$.*

Proof. For a chosen c and a proposed $b > c$, intersect the decision quorum at c with the promise quorum at b . A common acceptor cannot have supplied a free promise at b , since it accepted c . Its reported vote is at least c by condition 3. The greatest reported ballot m therefore satisfies $c \leq m < b$. If $m = c$, condition 4 gives the conclusion. Otherwise apply well-founded induction to the proposal at m , which exists by conditions 3 and 5. Finally, a chosen ballot has a proposal

by conditions 5 and 6 and decision-quorum nonemptiness. Trichotomy reduces agreement to the preceding argument or equality of ballots. $\square$

The checked declaration is `Synod.theorem8`; its intermediate induction lemma is `Synod.lemma7`. This proof follows UPaxos Appendix B.

B. Era-indexed families

Let $e(b)$ be a ballot's era and $e(i)$ the era of instance i . Phase I uses $\mathcal{Q}_{e(b)}^I$ and phase II uses $\mathcal{Q}_{e(i)}^{II}$. The configuration condition is

$$\forall e, \quad \mathcal{Q}_e^{II} \cap \mathcal{Q}_e^I \quad \wedge \quad \mathcal{Q}_e^I \cap \mathcal{Q}_{e+1}^{II}. \quad (1)$$

The reduction also requires proposal and choice era bounds:

$$\text{proposed}_i(b) \implies e(b) \leq e(i), \quad (2)$$

$$\text{chosen}_i(b) \implies e(i) \leq e(b) + 1. \quad (3)$$

Ballot order must respect era order. These conditions are separate from Equation (1); the latter alone is not a complete protocol contract.

Lemma 2 (Relevant intersection). *If $b_2 < b_1$, b_1 is proposed at i , and b_2 is chosen at i , then $\mathcal{Q}_{e(b_1)}^I \cap \mathcal{Q}_{e(i)}^{II}$.*

Proof. The bounds give $e(b_1) \leq e(i) \leq e(b_2) + 1 \leq e(b_1) + 1$. Thus $e(i)$ is either $e(b_1)$ or $e(b_1) + 1$. The first case uses the intra-era half of Equation (1), by symmetry; the second uses its cross-era half. $\square$

Instantiating the Synod theorem for every instance yields `Paxos.theorem10`. Its type also contains the promise, value-selection, acceptance, and choice conditions. Agreement holds for every history satisfying them, across any number of eras. This quantification does not establish that an operational sequence of membership changes preserves them.

C. Deriving the proposal-era bound

The inherited `EraLe` hypothesis, Equation (2), bypassed part of UPaxos Figure 2. The extension `MultiPromise.History` restores a known-instance frontier $i_{\max}$ and suffix promises $p_{\geq j}(a, b)$. A suffix certificate expands to a free promise at every $i \geq j$.

Theorem 3 (Suffix-promise reduction). *Suppose instance eras are nondecreasing. A promise at j , including a suffix promise, satisfies $e(b) \leq e(\min(j, i_{\max}))$. Each proposal at i has a nonempty phase-I quorum whose members supply a free promise at i , a last-vote report at i , or a suffix promise starting at some $j \leq i$. Then $\text{proposed}_i(b) \implies e(b) \leq e(i)$.*

Proof. Choose a member of the proposal quorum. For a point promise, $\min(i, i_{\max}) \leq i$. For a suffix promise, $\min(j, i_{\max}) \leq j \leq i$. In either case monotonicity of instance eras transports the promise's era bound to $e(i)$. $\square$

`MultiPromise.History.expanded_inv` additionally transports the full promise history to the inherited invariant structure. The resulting agreement theorem retains Figure 2's known-frontier bound for choices. The expansion is an

invariant reduction; the append-only evolution of configuration knowledge and the creation of suffix certificates still need an operational proof. A single certificate can denote every slot $j + d$ for $d \in \mathbb{N}$, as checked by `suffix_unbounded`. This proves representational unboundedness, not completion of those slots.

IV. GENERAL WEIGHTED INTERSECTION

For a finite support list L and natural weights w , let $W = \sum_{a \in L} w(a)$ and $W_q = \sum_{a \in L \cap q} w(a)$. Define $M_L(w) = \{q : 2W_q > W\}$. A duplicate-free list gives the usual finite-set interpretation; repeated entries instead contribute repeated weight. No bound on the cardinality of L is fixed in the theorem.

Lemma 4 (Disjoint selection bound). *For disjoint q, r and natural weights x, y ,*

$$2X_q + 2Y_r \leq X + Y + \sum_{a \in L} |x(a) - y(a)|.$$

Proof. At each node, at most one of the two selections contributes. The required pointwise inequality follows from $2x \leq x + y + |x - y|$ or its symmetric form; if neither selects the node it is immediate. Sum by induction over L . The Lean encoding uses $(x - y) + (y - x)$ with natural truncated subtraction. $\square$

Theorem 5 (Scaled weighted-majority overlap). *For positive integers k, l , if*

$$D = \sum_{a \in L} |kw(a) - lw'(a)| \leq 1,$$

then $M_L(w) \cap M_L(w') \neq \emptyset$.

Proof. Positive rescaling preserves strict majority. If two majorities were disjoint, their scaled sums would exceed their respective scaled totals by at least one each, since all sums are integers. Their combined excess is at least two, contradicting Lemma 4 and $D \leq 1$. $\square$

This is UPaxos Appendix A's Lemma 2, proved here by a pointwise discrepancy bound. `WeightedGeneral.scaled_overlap` is the checked declaration. The unit-change result, equal-rescaling result, and self-intersection follow as corollaries. Zero total weight yields an empty majority family, so self-intersection by itself does not assert availability.

The uniform bound is sharp. With two nodes, weights $(1, 0)$ and $(0, 1)$ have distance two and admit disjoint singleton majorities. This is checked by `distance_two_counterexample`. Some distance-two changes nevertheless preserve overlap; for example $(2, 2, 2, 0)$ to $(4, 2, 2, 0)$ does. Thus sharpness means that a larger universal allowance would admit counterexamples, not that all changes beyond the allowance are unsafe.

The inherited four-node weighted schedule remains an independent executable example. Its seven rows obey the required overlaps. It is useful for a concrete deployment shape; it is no longer the only evidence for the general weighted construction.

V. OPERATIONAL PROMISES AND FAULT CONTROLS

A. A checked acceptor transition system

An acceptor state is (f, H, R) : a promise watermark $f \in \mathbb{N}$, a newest-first vote history H , and a reply history R . Initially these are $(0, [], [])$. Acceptance requires $f \leq b$; a promise requires $f < b$. The transitions are:

$$\begin{aligned} \text{accept}_b(f, H, R) &= (b, b :: H, R), \\ \text{promise}_b(f, H, R) &= (b, H, (b, \text{head? } H) :: R). \end{aligned}$$

An empty head produces a free reply; otherwise the reply carries the latest accepted ballot. The histories are ghost evidence for the proof. This module does not prescribe retaining an unbounded physical list at a replica.

Theorem 6 (Historical promise safety). *In every state reachable by a finite execution, a free reply at b excludes votes below b . A reply carrying c at b establishes $c < b$, an acceptance of c , and absence of acceptances strictly between c and b .*

Proof. Strengthen the induction invariant: all votes are at most the watermark; the head vote dominates its tail; every emitted reply's ballot is at most the watermark; and each reply satisfies the stated historical conditions. An acceptance preserves every old reply because its ballot is at least the old watermark. A promise reads the maximal vote and raises the watermark above it. Neither transition removes prior evidence. Induction on reachable states proves the claim. $\square$

The declarations are `Acceptor.reachable_inv`, `free_forbids`, and `report_last`. They discharge S2 and S3 for a single-era acceptor. They do not establish the proposal value rule, an entire log's view-change selection, or restoration after volatile-state loss. Extending the component to era-tagged views also remains an explicit integration task.

B. What the negative controls establish

The controls separate two questions: whether the proof notices an altered definition, and whether removing an assumption admits an inconsistent history. The first is measured by compiler rejection; the second by constructing and checking a counterexample with the remaining conditions intact.

- *Cross-era overlap.* The inherited two-node history satisfies intra-era overlap and the other encoded era-history conditions, but chooses both Boolean values at one slot when the cross overlap is absent.
- *Value selection.* One acceptor reports ballot 0 at ballot 1, yet the two proposals use different values. All listed Synod conditions except S4 hold and both ballots are chosen.
- *Promise fencing.* One acceptor supplies free promises at two rounds despite its earlier vote. All listed conditions except S2 hold, and agreement fails.
- *Nonempty decisions.* Empty phase-II quorums permit choices with no proposal or acceptance. The six history invariants hold vacuously where appropriate, and the two selected values differ.

- *Operational guard.* After a free promise at round 2, accepting round 1 without checking the watermark violates the reachable invariant.

These examples establish independence within the specified contracts. They do not prove a universal necessity theorem over all consensus algorithms, nor do they enumerate every combination of omitted assumptions.

Separately, the mutation harness replaces the checker’s universal quantifier over quorums by an existential one, and replaces a safe concrete quorum family by an unsafe family. Lean rejects each altered source while accepting the unchanged control. An injected `sorry` demonstrates a different failure mode: compilation succeeds but `#print axioms` exposes `sorryAx`. The harness bounds child memory, work and wall time and modifies only temporary copies. Infrastructure failures are not counted as killed mutants.

VI. THE BOUNDARY BETWEEN UPAXOS AND VRR

VRR-2012 is a replication protocol, with client handling and recovery in addition to the agreement mechanism [4]. Its normal path requires ordered log preparation and a quorum of acknowledgements before execution. View change selects the log with the greatest last-normal view, breaking ties by length. That whole-log selection rule needs a proof relating logs from different views; a per-slot maximal-vote rule is not automatically the same operation.

The distinction becomes essential after crashes. A recovering replica is excluded from normal processing and view changes. Its recovery exchange requires responses from distinct peers, a fresh nonce, and state from the primary of the greatest reported view. The preceding round of `StartViewChange` messages provides replicated knowledge of view advancement. The paper’s Section 8.2 gives a failure scenario when that round is removed. Replacing volatile state by an acceptor watermark that never disappears would bypass precisely this obligation.

The 2012 reconfiguration protocol stops taking new client requests in the old epoch, transfers the committed state, and activates the replacement group. Section 7.5 explicitly discusses the resulting delay and prewarming as a way to reduce it. The `uVRR` objective instead calls for progress during the transition. Equation (1) and casting-vote scheduling address part of that difference, while the interaction with recovery remains to be proved.

A. Checked whole-log selection

The executable `ViewSelection.select` takes a seed report and a list of additional reports, enforcing nonemptiness without a fallback value. It selects by last-normal view and then log length. Each report’s committed frontier is bounded by its length in the report type. The checked results establish that the output belongs to the input collection and has maximal rank. For reports of length at most M , `scalar_rank_equivalent` proves that the model’s scalar rank $v(M + 1) + n$ has exactly the same ordering. A counterexample shows that omitting the length bound can misrank reports.

The prefix-preservation result is the induction step for a single era. Let a prefix P have committed in view v . Intersect its commit quorum with the reporting quorum and obtain a report from a common sender. The explicit voter-history premise says that this report has retained view at least v and, at equality, contains P . The induction hypothesis supplies preservation for reports from strictly later normal views. If the selected report is later than v , that hypothesis applies. Otherwise the witness and selected report are both from v . Same-view log comparability and maximal selected length then imply that the witness log, and hence P , is a prefix of the selected log. `ViewSelection.quorum_preserves` checks this argument.

The selector itself only compares ranks. The surviving-voter and log-origin premises still require proofs from message and recovery transitions. A two-report fault control admits divergent uncommitted suffixes in different views: an older two-entry log loses to a newer one-entry log. A length-only selector chooses the older log and loses the newer committed entry. These results close the selection subproblem while preserving the distinction between a checked induction step and a complete reachable-state proof.

B. Normal-log message provenance

`NormalLog` first checks the fixed-view message argument. The primary starts from an arbitrary base and appends. `StartView` installs that base once; normal replicas accept authentic prepares at their next slot and emit reports copying their logs. Persistent messages permit delayed, reordered, and duplicate delivery attempts. The invariant makes replica logs, sent prepare prefixes, and reports prefixes of the primary history. At receipt, the replica and pre-send logs are equal-length prefixes of that history, hence equal; appending the sent entry preserves the invariant. The pre-send log is ghost evidence, not a content check in the handler. `reports_comparable` follows from prefix comparability, and `replica_history` proves retention over finite continuations. Bypassing next-slot checking admits an authentic second entry at slot one. This component supplies the local argument integrated below; its arbitrary initial base alone does not justify cross-view safety.

C. Shared multi-view preservation

The next integration is now checked. `ViewFence` models a local replica installing higher normal views, appending in normal mode, raising its floor before reporting, and retaining immutable vote and reply histories. Its `report_covers_vote` theorem bounds every historical vote below a report’s target by the reported last-normal view; at equality, the voted log is a prefix of the report. This remains true of votes later in the trace. The selector therefore consumes a derived voter-history property. Strong induction discharges the strictly-later-view premise over activated views; a concrete two-view execution demonstrates that the interface is inhabited.

`LogProvenance` joins the local histories to a shared message model. Each view has one immutable installation base and an append-only primary history. Activation requires

a quorum of authentic reports and uses the executable latest-normal-view/length selector. An issued prepare carries a view, next slot, and entry; its pre-send log is ghost evidence. Installation reads an activated base. Receipt checks normal mode, view equality, and the next slot. Fencing and report emission update the local replica state. There is no guard testing global committed-prefix agreement.

The invariant establishes that each replica log, historical vote, and report extends its retained view’s base and is a prefix of that view’s primary history. Each activated base retains its selection certificate because report histories persist. These facts discharge both same-view comparability and base-extension provenance in the view induction. `later_base_preserves` proves preservation by every later activated base. `committed_comparable` proves that any two quorum-voted prefixes in a reachable state are comparable; `committed_equal` specializes this to equal-length prefix equality. The model bounds neither log length, view numbers, nor activation count, and fixes one self-overlapping quorum family. A concrete trace commits a nonempty prefix in two successive views. A delayed old-view prepare has the expected next slot after installation; bypassing view equality permits an unsupported new-view vote. The constructive control and source mutation both check this boundary.

This is a complete safety induction for the stated crash-free shared model, not yet the requested uVRR theorem. Unique source activation is enforced by an unused-view guard, and source histories may continue issuing prepares after replicas have fenced that view. Their eventual acceptance is constrained by replica guards. A refinement must justify these abstractions from unique primary activation and the concrete wire protocol. Crashes, first-round view-change knowledge, arbitrary membership changes, client responses, and progress remain outside this model. Ghost histories do not imply durable storage at replicas, but they do not prove recovery after volatile-state loss.

D. Recovery fencing: a checked boundary

`RecoveryFence` proves a narrower persistence result. Suppose a fixed support quorum is simultaneously online with bounds at least b , and historical responses carry causal recipient-episode identifiers. Crashing erases a local bound and advances the environment’s fresh episode identifier; recovering nodes cannot answer. Recovery uses distinct current-episode responders and chooses a maximal bound. Induction proves that every online member of the support retains a bound at least b through arbitrary finite crash/recovery continuations. The quorum checkpoint is a premise, not a consequence of collecting asynchronous view-change messages. The model does not restore logs or establish progress. A stale-episode witness and a source mutation check the freshness boundary.

Classic diskless crash-recovery designs are exposed to a known failure pattern: Michael et al. describe delayed view-change and recovery divergence in Appendix B.1 and Figure 1 of their extended report [1], [2]. That is literature about the classic crash-recover class, not an open bug here: uVRR eliminates the class by construction. A crashed node whose

durable superblocks record an unflushed session reopens under a new incarnation (Crash-Stop-Self-Evict), the leader evicts the old identity, and the new identity rejoins as a weight-0 learner that may receive streams but whose messages are ignored. Same-identity recovery after volatile-state loss is therefore unrepresentable, and the stale-episode rule above is precisely the ingress rule the reincarnated learner obeys. Their crash-vector construction provides the collection primitive the reincarnation leader needs.

E. Crash-vector collection

Following Algorithm 1 and Definition 6 of the extended report [2], let a reply r contain sender s_r , vector $v_r : A \rightarrow \mathbb{N}$, and a payload. The collector stores a known vector k and replies R . Upon receiving r , it computes

$$k'(a) = \max(k(a), v_r(a)), \\ R' = \{q \in r :: R \mid k'(s_q) \leq v_q(s_q)\}.$$

The implementation uses a list; quorum membership is the set of sender identities. Thus duplicates do not create additional identities. Induction proves every retained vector is pointwise bounded by k . The filter gives the converse inequality at each retained reply’s own sender. Consequently,

$$\forall r, q \in R, \quad v_r(s_q) \leq v_q(s_q),$$

the published crash-consistency condition. A known-stale reply remains ineligible as the frontier grows. A concrete three-node example shows that an unfiltered two-sender quorum violates this condition; filtering removes the stale reply and a current replacement restores collector quorum size. Removing the filter in a temporary source copy causes Lean rejection.

This is a collector theorem, not a recovery theorem. The published protocol requires recovery itself to acquire and propagate its new incarnation through a crash-consistent quorum; in uVRR that acquisition is the reincarnation path the four-superblock identity contract drives. Its quorum-knowledge persistence argument also covers reconstruction, message freshness and resending after pruning. Those obligations are not discharged by attaching vectors to ordinary VRR replies. The report’s termination condition requires an appropriate quorum stable over the acquisition interval; arbitrary churn with a bound on simultaneous recoveries alone is not a termination theorem. Mechanizing that temporal argument and composing it with log safety remain open.

F. Acquisition ordering and operational certificates

`AcquisitionOrder` checks the secondary induction of Theorem 3 in the extended report [2]. With asynchronous participant witness times and a generic stable property, an on-line participant either retains its witness knowledge or inherits it from an earlier response quorum. Quorum intersection and forward ordering of those responses relative to the witness events discharge a strong induction over time, using core Lean’s `Nat.strongRecOn`. A separate lemma derives that forward ordering from crash-vector consistency, conditional on recovery participants retaining incarnation knowledge. A countermodel satisfies every other origin condition except

forward ordering and loses the property; a source mutation that removes the ordering premise is rejected.

`RecoveryAcquire` then makes the acquisition operational, following Algorithm 1. Crash, start, emit, answer, collect and finish transitions act on an immutable authenticated message history. A crash erases all protocol knowledge and begins a fresh logical incarnation; request identity is an explicit echoed (incarnation, sequence) pair in addition to the response vector check. Induction over reachable states proves that every finished certificate is a crash-consistent quorum for exactly its request and incarnation, and that every historical reply originated at an operational sender. A three-node trace completes recovery at incarnation 1. An authentic reply from the same incarnation but an older request violates freshness when replayed into a new acquisition, which separates sequence freshness from incarnation freshness. Mutations that answer from a recovering sender or accept a stale request are rejected.

These are acquisition mechanics with one active acquisition per node. Quorum-knowledge persistence across the acquisition interval, value reconstruction, the read quorum, mapping the echoed (incarnation, sequence) identity to the superblock incarnation contract, log integration and liveness remain open.

G. Unbounded VRR reincarnation

uVRR does not implement crash-recovery. It enforces Crash-Stop-Self-Evict, called *reincarnation*: a node whose volatile state was lost reopens as a different node. Durable state is four TigerBeetle-style superblocks. A restart that reads all four as `flushed` marks them `unflushed` and continues normally; a restart that reads any `unflushed` block is dirty, bumps its incarnation, and rewrites all four blocks as the new identity marked `flushed`. Any read that observes a higher identity adopts it. Once the eviction is initiated it must continue. A *reincarnation* message carries the identity pair (old, new) to the leader; the durable incarnation supersedes the fence machinery’s ghost generation as the freshness carrier where the two meet, while the fence and acquisition theorems above remain in force with the carrier re-instantiated.

The leader stays live for client requests and drives a forced weight sequence from the voting-weights rules [7]: the exiting node’s weight goes $1 \rightarrow 0$, the new identity joins at 0, and the new identity goes $0 \rightarrow 1$. Each step is a unit-weight change, so the strict majorities of consecutive configurations overlap by Theorem 5’s corollaries and every intermediate era is quorum-safe on its own; a leader dying mid-sequence leaves only safe eras that a new leader can continue. If the sequence runs in the doubled scale of a weight-doubling plan, the return halve is integral only when all weights are even, so the corner sequence inserts one extra unit increment on the new identity before the halve. The leader streams the first reconfiguration, then client traffic, and preemptively streams to the reincarnated node as a weight-0 learner; messages to it are fine and messages from it are discarded by the standard membership check, which is the same rule the crash-vector collector’s stale-reply filter realizes. If the leader crashes, the cluster must reach a stable leader before the reincarnated node forces its old identity’s eviction.

The economics defer disk flushes: nano-state fits in memory, so disk flushes can be deferred, and the network round trips

of reincarnation are cheaper than fsyncs because “the network is faster than the disk” [6]; the four-superblock mark is one fsync amortized over the whole epoch. Same-identity recovery after volatile-state loss is thereby unrepresentable rather than repaired: classic VRR-2012 diskless quorum recovery is not performed, so the DISC’17 Appendix B.1 amnesia class is avoided by construction, and a zero-weight learner may speculatively recover its log non-votingly as future work. The formalization of the superblock contract and the forced sequence against the fence and acquisition machinery is a concrete proof obligation, planned as a later ladder rung.

H. Casting votes and the design model

The inherited casting-vote lemma assumes an old decision quorum q , a new promise quorum q' , and a leader ℓ with $q \cap q' = \{\ell\}$. Sending new promises only to $q' \setminus \{\ell\}$ leaves the old acceptance guard unaffected at the other members of q . Once ℓ promises, the new quorum is complete. This is a schedule-dependent noninterference argument. Availability of those quorums, continued processing by the leader, successful message delivery, and the ability to repeat the procedure require additional proofs. The lemma includes no clock or duration.

The repository’s `VrrCoreEras.tla` is explicitly a design model for a single era transition, rather than a statement that the Rust implementation already refines it. Its gate checks both cross-era view/commit directions and fence intersections. The Lean history theorem’s narrower directional overlap must not be used to remove those checks without proving the relevant transition refinement. Table I records the outstanding source-level obligations, including optional VRR optimizations whose semantics differ from the base protocol.

VII. REPRODUCIBILITY AND RESEARCH METHODOLOGY

A. Executable evidence

Each positive rung contains the entire module text, a build invocation, and axiom reports for its stated results. `showboat verify` reruns the commands and compares their captured output [8]. All eight inherited documents pass this check. The eighth contains the expected compiler failure of a historical Leanstrat draft, so its successful replay does not establish the missing weighted theorem. Rung 9 supplies that theorem.

The checked additions are numbered 9 (general weights), 10 (operational acceptor), 11 (suffix promises), 12 (fault controls), 13 (whole-log selection), 14 (normal-log provenance), 15 (fenced histories), and 16 (shared multi-view preservation), 17 (conditional recovery fencing), 18 (crash-vector collection), 19 (acquisition ordering), and 20 (operational acquisition certificates). The classic crash-recover attempt’s delayed-fence counterexample rung was deleted with that attempt, and the ladder’s recovery rungs are numbered 17–20 after renumbering. The original report is retained as an audit input. Its statement that every rung includes a successful build and axiom report does not hold for rung 8.

The independent baseline audit ran the unmodified library build and every original transcript, inspected theorem types and axiom reports, and preserved input hashes. Subsequent

TABLE I
COVERAGE OF THE WHOLE VRR-2012 PAPER AT THIS CHECKPOINT. “OPEN” MEANS NO END-TO-END LEAN RESULT IS CLAIMED FOR THAT OBLIGATION.

Source section	Required obligation	Checked evidence and present boundary
1–3: model and architecture	Crash-only failures, authenticated identities, network model, deterministic execution, replica thresholds	Explicit assumptions; quorum intersection is checked. No general environmental refinement theorem.
4.1: normal operation	One value per view/slot, ordered preparation, distinct quorum support, commit before execution, request deduplication	Conditional slot agreement; operational acceptor promise rules and fixed-view log provenance. Proposer, clients, and execution integration open.
4.2: view changes	Sender status, two message rounds, latest-normal-view/length selection, preservation of every committed prefix	Shared crash-free model proves whole-log compatibility across unbounded views. Concrete protocol refinement and the reincarnation proof (the Crash-Stop-Self-Evict forced weight sequence) remain open.
4.3: replica recovery	Exclusion until fresh quorum evidence and current-primary state; restored votes consistent with prior activity	Conditional fence retention and crash-vector collection checked; the classic crash-recover class eliminated by construction (Crash-Stop-Self-Evict reincarnation).
4.4–4.5: applications and clients	Chosen nondeterministic inputs; monotone request identifiers across client crashes	Explicit application/client refinement obligations, open.
5.1–5.3: practical transfer	Checkpoint fidelity, suffix completeness, garbage collection, view-sensitive truncation, short view-change messages	Open; source paper’s full-log argument does not validate optimized transfer automatically.
6.1–6.2: witnesses and batching	Witness participation roles; batch semantics and operation order	Optional refinements open. No correctness claim inherited merely from batching.
6.3: fast reads	Lease exclusion and clock assumptions, or the weaker semantics of backup reads	Open; the backup-read variant is not claimed linearizable by the source.
7.1–7.5: reconfiguration	Epoch establishment, transfer completion, safe retirement, duplicate messages, client routing	Era-history reduction and quorum calculations checked; operational repeated uVRR reconfiguration open.
8.1–8.3: correctness arguments	Compose normal operation, recovery and reconfiguration; state progress assumptions	Component safety results checked. Composition and liveness open.
9: conclusion	Relate the completed results to the replication service	No additional theorem is inferred from the source’s summary.

changes were checked in small increments. The library build and transcript checks complement the Rust format, lint, test and documentation gates. Passing the Rust suite does not establish a formal refinement to Lean. A vendor-directory match in an inherited repository text gate is documented rather than suppressed by modifying unrelated code.

An independent TLC 2.19 breadth-first run of `VrrCoreEras.cfg` completed in 137.67 seconds, generating 2,743,933 states and finding 837,204 distinct states, with no error reported and an empty remaining queue. Its scope is the three-node `inc3` scenario, one command value, maximum log length three, two eras, and view index zero. `MaxEpoch=0` disables crashes; this run supplies no crash-recovery coverage. TLC reported a fingerprint collision probability estimate of 2.2×10^{-7} from the observed fingerprints. This finite exploration is separate from kernel-checked unbounded theorems. The `VrrCoreErasM1.cfg` transfer mutation produced the expected `FrontiersOrdered` violation in 1.2 seconds. Source and configuration hashes and full logs are retained for both runs.

A single executable document, `REPRODUCE.md`, packages the whole chain: the exact tool versions (Lean 4.33.1, Lake 5.0.0, Showboat 0.6.1, TLC 2.19 from the `v1.7.4` release

of `tla2tools.jar` pinned by SHA-256, Tectonic 0.17.0 and the Rust 1.96.0 toolchain), the library build, the axiom audit, every rung replay, the mutation controls, the two TLC runs, the Rust gates, the paper build and SHA-256 digests of every input. `showboat verify REPRODUCE.md` reruns all of it and diffs the captured output. The paper build fixes `SOURCE_DATE_EPOCH`, so the rendered PDF is byte-identical across builds with the same Tectonic release and package bundle, and its digest is recorded. The mutated-configuration TLC run is replayed with one worker, because a multi-worker search stops at a nondeterministic point after the first violation; only the violated invariant is compared.

B. Bounded multi-model proof assistance

Leanstrat is a Lean-oriented model available through Mistral’s tooling. Its release article reports benchmarks and case studies [9]; those results do not establish cost effectiveness for this core-only development. Here the model is a source of candidate proof text. The local Lean kernel, the theorem type, and an explicit axiom policy determine acceptance.

The inherited record reports a failed attempt on the general weighted lemma, including a large Lean process. We preserved that record and replayed the draft’s compiler failure. We did

not independently reproduce its historical peak memory or elapsed time.

The new experiment isolated a single list-induction obligation with unchanged definitions and statement. The first Vibe run had a 12-turn, 28,000-token, 180-second and 2-GiB sampled process-tree memory limit, with a configured USD 2 cap. It stopped after 8.19 seconds when scaffold usage exceeded the token limit, before a proof edit. The second used 6 turns, 150,000 tokens, 120 seconds and a configured USD 0.50 cap. It stopped after 76.39 seconds on the cumulative-token limit, leaving the original proof hole. Peak sampled resident memory was about 205 MiB in each run. Neither cap is a measurement of actual billing, and sampling is not an operating-system memory guarantee. Two direct-API routing probes also failed. The weighted theorem was then completed and compiled locally.

No accepted Leanstrat-generated proof is attributed to these runs. Their methodological result is that scaffold overhead dominated a tiny proof target; repeating the same configuration was not justified. Future comparisons should record exact prompts, tool context, candidate edits, compiler iterations, token usage where available, elapsed time, and invoiced cost. A claim of multi-model efficiency needs comparable successful obligations and a baseline workflow. This experiment supplies neither a measured dollar saving nor a general negative judgment about the model.

The laboratory book distinguishes user-supplied authorship of earlier work from independently checked artifacts. This permits credit to prior model assistance without making model identity part of the trusted proof base.

VIII. NEXT PROOF BOUNDARY AND PERFORMANCE CLAIMS

The shared crash-free model now preserves quorum-voted prefixes across view changes. The next integration theorem must relate its activation, message-history, and local-state abstractions to the concrete VRR protocol. The Crash-Stop-Self-Evict reincarnation path (the four-superblock durable identity, the reincarnation message, and the forced weight sequence) must be proved against the fence and acquisition machinery checked here. Finally, reconfiguration must preserve these facts for an arbitrary sequence of membership changes and state transfers. These are concrete proof obligations, not consequences of quorum intersection alone.

The term *non-stop* also needs an operational definition. A useful conditional progress statement would quantify over an available casting-vote schedule and fair delivery while allowing arbitrary pipeline length. It would state which classes of requests keep completing while the new phase-I exchange is pending. It would not promise service after loss of all necessary quorums. Similarly, no fixed failover time follows in an asynchronous model allowing unbounded message delay. A single-digit-millisecond result requires specified failure detection, network and workload conditions, an implementation, and measured latency distributions.

The present results establish exact component theorems with executable evidence and targeted counterexamples. Completing the operational composition will determine the scope of

the final uVRR correctness claim. Performance and novelty claims will be evaluated against that completed specification and its experimental evidence.

APPENDIX REPRODUCTION COMMANDS

From the repository root, with the tools listed in `formal/uvrr-lean/REPRODUCE.md` installed at the recorded versions:

```
export PATH="$HOME/.elan/bin:$PATH"
cd formal/uvrr-lean
showboat verify REPRODUCE.md
```

That document lists every command with its captured output, and `showboat extract REPRODUCE.md` emits a shell script that recreates it. The document runs the pinned Lean toolchain declared in `lean-toolchain`. Showboat transcripts compare command output, not only exit status, so a fresh library build is recorded independently of transcript replay. Historical rejected drafts are quarantined outside the compiled library.

REFERENCES

- [1] E. Michael, D. R. K. Ports, N. K. Sharma, and A. Szekeres, "Providing Stable Storage for the Diskless Crash-Recovery Failure Model," UW-CSE-16-08-02, 25 August 2016. <https://homes.cs.washington.edu/~drkp/papers/diskless-tr16.pdf>.
- [2] E. Michael, D. R. K. Ports, N. K. Sharma, and A. Szekeres, "Recovering Shared Objects Without Stable Storage [Extended Version]," UW-CSE-17-08-01, 5 August 2017; extended version of DISC 2017 paper. <https://drkp.net/papers/recovery-tr17.pdf>.
- [3] D. C. Turner, "Unbounded Pipelining in Dynamically Reconfigurable Paxos Clusters," revision 1A9DBA37, 14 August 2017. <https://tessanddave.com/paxos-reconf-latest.pdf>.
- [4] B. Liskov and J. Cowling, "Viewstamped Replication Revisited," MIT-CSAIL-TR-2012-021, 23 July 2012. <https://hdl.handle.net/1721.1/71763>.
- [5] simbo1905, "Unbounded Viewstamped Replication Revisited?" 12 August 2026. <https://simbo1905.wordpress.com/2026/08/12/viewstamped-replication-revisited/>.
- [6] simbo1905, "The network is faster than the disk," 12 April 2024, accessed 7 September 2026. <https://simbo1905.wordpress.com/2024/04/12/the-network-is-faster-than-the-disk/>.
- [7] simbo1905, "Paxos Voting Weights," 16 March 2017, accessed 7 September 2026. <https://simbo1905.wordpress.com/2017/03/16/paxos-voting-weights/>.
- [8] S. Willison, "Showboat: Create executable demo documents," software, version 0.6.1 used for the inherited evidence. <https://github.com/simonw/showboat>.
- [9] Mistral AI, "Leanstrat: Open-Source foundation for trustworthy vibe-coding," 16 March 2026, accessed 6 September 2026. <https://mistral.ai/news/leanstrat/>.